




NOSQL With Grails

Joseph Nusairat
Groovy Sage
@nusairat

About Me

- * Java Developer since 1997
- * Groovy / Grails Developer since 2007
- * Scottsdale Groovy Brigade co-founder



Abstract

- * What is NoSQL
- * How to use NoSQL
- * How to call MongoDB
 - * How to set up Mongo Documents
 - * Querying Mongo
- * Using Mongo with Grails
- * Examples



NoSql

- * Term originally coined in 1998 by Carlo Strozzi
- * Reintroduced in 2009 by Johan Oskarsson
- * Is used to refer to databases that do not use the standard SQL interface



NoSql Types

- * Column: HBase, Accumulo
- * Document: MarkLogic, MongoDB, Couchbase
- * Key-value: Dynamo, Riak, Redis, Cache, Voldemort
- * Graph: Neo4J, Allegro, Virtuoso



Design

- * Transactions are fast because there are none. No overhead of transaction managers
- * There are no tables or joins in the traditional sense. Meaning access to multiple tables with one query is not possible



In Memory

- * Part of the speed performance is that NoSql often uses in memory storage
- * This is generally a temporary store till it flushes to the file system, but this can lead to volatility if the server goes down.



Using NoSQL

General Use

- * RDBS systems and traditional database design in mid to late 2000s started to show their age
- * Many developers started switching ad hoc to NoSQL systems to speed up the applications
- * Became the defacto toy to the shiny object crowd



Issues

- * Sacrificed normalized design for performance
- * No transactions meaning you'd have to make sure you were not doing multiple saves
- * Potential for over duplication of data





Speed vs Design



Not Only SQL

- * A more accurate representation of the name in how to use it
- * NoSQL is used for high performance and throughput of data
- * Any part of your site that gets used by spiders or heavy searching by people
- * Increase performance on what needs performance



Where to Use

- * Front end facing features
- * Pages linked on the front end
- * If its a a sales site, the items you are selling
- * Huge non transactional calls
- * Huge amounts of data



When Not to Use

- * Highly complex structured relations that you want to preserve or change
- * Inheritance where updating a sub item effects others
- * Rarely used data
- * Small data sets



Combination

- * Sometimes you have complex data that needs to be accessed quickly
- * Data that needs back end complexity to set up and manage, but needs to be easily accessed on the front end
- * Duplication is O K.



Example

* Deal Site

- * A user creates a complex deal with locations, versions of the deal, multiple deals
- * Store the inheritance for the user but then create a NoSQL record for the front end to run and query on



Threshold?

- * Depends on hardware
- * Depends on database size
- * Also consider with document based NoSQL the taxing of the database is much less



Possible Exceptions

- * Multiple database servers can require multiple expertise and great amounts of money
- * If you have only a few collections / tables that you'd want to store in RDBMS then it may not make sense from a cost view



Mongo DB

What is it

- * Document type database originally written in C++
- * Data is stored in Binary JSON (BSON) which allows for binary data type
- * Data is arranged in collections which also allows for collections embedded in collections



What Makes it Special

- * Document databases allow for more complex storage of data in an easy to read way (reads like JSON)
- * Can also store binary files on the file system with a reference to them
- * Allows for location based queries



Everything's Dynamic

- * Traditional databases you define your databases, tables, etc ahead of time
- * With Mongo it can all be defined on the fly
 - * Databases
 - * Collections



Collections

- * Collections will be stored with BSON
- * Collection by default contains a unique “_id”
- * Collections can contain embedded documents or even lists of embedded documents
- * Collections do not have a distinct set of columns
- * Collections store documents



A Document

An Entire Collection

```
{  
  "username" : "bob",  
  "street" : "123 Main Street",  
  "city" : "Springfield",  
  "state" : "NY"  
}
```



Collections in Documents

- * When you need to reference collections from a document you can handle in 2 ways
- * Embedded - having all the objects in the collection
- * Reference - having the documents in one collection then referencing them in another



Embedded

- * You can embed documents either in a list format or just one document at a time

```
{  
  "username" : "bob",  
  "address" : {  
    "street" : "123 Main Street",  
    "city" : "Springfield",  
    "state" : "NY"  
  }  
}
```

Contains only one



Embedded

- * A list will look the same but have an [] around the items.

```
{
  "username" : "bob",
  "address" : [ {
    "street" : "123 Main Street"
  },
  {
    "street" : "234 High Street"
  } ]
}
```

Contains many addresses



Reference

- * You can also reference the collections from another collection, this will look more like a SQL type collection.
- * For collections with data stored outside
- * For extremely large collections
- * Rarely used data



Reference

- * References can be referenced either with the `_id` directly or with `DBRef` object

```
{
  "_id" : ObjectId("52a848d8 ... b0b3"),
  "name" : "My Item 1",
  "locations" : [
    DBRef("locations", ObjectId("5"))
  ]
}
```



Inserting a Document

- * If you want to insert the previous document into a collection called “people”

```
db.people.insert({  
    "firstName" : "joseph",  
    "lastName"  : "nusairat",  
    "address"   : {  
        "street" : "123 Main Street",  
        "state"  : "NY"  
    }  
})
```



Inserting a Document

- *The result in the database is creation of an object with the id “_id”

```
{  
  "_id" :  
    ObjectId( "4cda8571b5da950b52727746" )  
  "firstName" : "joseph",  
  "lastName" : "nusairat",  
  "address" : {  
    "street" : "123 Main Street",  
    "state" : "NY"  
  }  
}
```



Updating Records

- * Record updates in Mongo is much different than in normal RDBS systems.
- * Default is to update entire record
- * Can do an update where it will insert if the record does not exist



Default Update

- * Default call will take in a set of parameters that will be the “find” part of the update
- * The second item is what to update
- * This will update people to having just the column first name

```
db.people.update(  
    { _id : ObjectId("4cda...7746") },  
    { firstName : 'joseph' }  
)
```



Setting one Column

- * If you only want to set one column you can use the \$set call
- * Using the set will update the column if it exists, or add it if it does not exist

```
db.people.update(  
  { **condition** },  
  { $set :  
    { firstName : 'joseph' }  
  }  
)
```



Other Choices

- *Upsert : You can choose whether to add a record if it doesn't exist
- *You can choose whether to update all records or the first match

```
db.people.update(  
  {}, { $set : {middleName : 'Javier'},  
  true, true } )
```

If all documents matched in the criteria are to get updated

Upsert : if the record does not exist we should insert it



Removing a Record

- * Removing record is a direct call
- * The call will remove all records that match a given set of indicators

```
db.people.remove({ firstName : 'joe' })
```

Defines the syntax to search for on removals.



File Storage - Grid FS

- * GridFS is used for storing files. This is used more specifically to store files that are over 4MB
- * This allows for safely storing large files by dividing them up among multiple documents
- * Stores the files in different buckets the default bucket is called “fs”



Example with Groovy

- * Becomes a bit more complex with direct calls
- * Returns an id for correlation

```
def gridfs = new GridFS(db)
def f = new File('/var/in/struts.png')
def inputFile = gridfs.createFile()
inputFile.save()
```



Querying Mongo

Querying

- * Querying records is more similar conceptually to those using GORM.
- * You can query and find one record or find all

```
db.people.find({ firstName : 'joe' })
```

Find all records that

```
db.people.findOne({ firstName : 'joe' })
```

Find the first match



Map - Reduce

- * Is used for more complex querying from the database
- * This helps with the aggregation of data
- * Map reduce can be used in a situation where you would normally have used group bys in SQL.
- * Written in JavaScript
- * Data outputted to temporary table



How it Works - Map

- * Mapping is designed to take a large input and then divide up into smaller pieces
- * The map explicitly defines the item that is going to be aggregated on



How it Works - Reduce

- * The reduce aggregates the map outputs.
- * Taking the smaller pieces and bringing them back into one for the final item
- * The aggregation is based on the key passed in



Example

```
// emit calls the next result
// first: differentiator / second: aggregator
function map() {
    emit(1, {count : this.amount});
},

function reduce(key, values) {
    var count = 0
    for (var i = 0; i < values.length; i++)
        count += values[i].count
    return {count: count}
}
```



Geo Spacial Querying

- * Has the ability to set up a per collection field that is geo spatial specific
- * These fields need to be defined with an index “2d”
- * Can query records based on location



Geo Querying Types

- * **Near** - find the items with the closest locations to you in order of closest first
- * **Box** - define boundaries of a box and find all matching items in that box
- * **Circle** - define boundaries of a circle and find all matching items in that circle



Mongo With Grails



Stand Alone or GORM



Mongo with Grails

- * Mongo can integrate with Grails either via Mongo API calls directly or by a GORM syntax
- * Both syntax's can be used in the same application
- * Will depend what you are comfortable with



Inline

- * Inline syntax will use Mongo directly
- * All creates / removes / queries will be with Mongo syntax and have to convert to local objects
- * The syntax this way can be more precise when creating complex queries



GORM

- * Makes use of the Grails sponsored plugin : mongodb
- * This plugin wraps a GORM syntax around calls to Mongo
- * Obscures the actual calls to Mongo
- * Allows for embedding or referencing collections and domains



Adding Mongo Gorm

- * Need to add the plugin
- * Define the data source for Mongo
- * If you are using a mix mode system must include:

```
static mapWith = "mongo"
```



Identifier

- *nothing

- *By default will use a sequence number storing a sequence in Mongo

- *I wouldn't recommend

- *String id

- *Sets a Uuid

- *ObjectId ***recommended***

- *Uses the Mongo Object id



Querying with GORM

- * Querying with Mongo uses the same Hibernate style syntax

- * There is custom syntax for the Geo Queries

 - `findByLocationNear`

 - `findByLocationWithinBox`

 - `findByLocationWithinCircle`



Issues With GORM

- * When changing the object on the embedded collection, GORM will not automatically update the embedded objects
- * Limited to geo querying on first levels of the collection
- * 2D on embedded lists



Examples

THANK YOU - Q&A

<https://github.com/nusairat/mongo-groovy-example>